Project Proposal
Restaraunt Segmentation Analysis

Capstone Partner: Sitewise Analytics
Contributors: Chen Lin, Eric Tsai, Morris Zhao, Xinru Lu
Mentor: Gittu George

2023-05-08

Contents

Executive Summary

The Restaurant Segmentation Analysis project aims to use machine learning methods to determine factors that drive traffic to a particular location and identify clusters of similar store locations. In the following proposal, we will first highlight the problem and fundamental goals of the project, then discuss the data science techniques that we will employ to tackle the problem, and conclude with a rough timeline of the project.

Introduction

Restaurants seeking to open new stores in a region need to make marketing plans according to the major customer group, whether it is office workers downtown or students in a neighbourhood. Therefore, restaurant franchise owners need to know the factors that drive traffic to a location, such as the surrounding population demographic and consumer behaviour in the region, as well as trade area and nearby competitor/sister store information. By having a strong grasp of these factors, owners can plan future expansions and market the new location strategically based on the demand of the region. Given that Sitewise Analytics does not currently have these factors, the Restaurant Segmentation Analysis project will address this problem by using data from Smoothie King locations in the United States and Subway locations in Canada and the United States to build machine learning data pipelines for Sitewise to incorporate into their consulting service for those respective clients. At the end of the project, we expect to have human-interpretable machine learning models that cluster similar store locations, which will be helpful for Sitewise clients to identify factors that drive traffic in those similar locations.

Given that Smoothie King, Subway US, and Subway Canada are all different clients of Sitewise, it is necessary to build three separate models for each respective client. Ultimately, the factors that drive traffic to each of the three restaurants may be different as well. Thus, the three machine learning pipelines are as follows:

1. A supervised machine learning pipeline using data from Smoothie King locations to predict a store's category from one of five categories:

   - Home
   - Shopping
   - Work
   - Travel
   - Other

The prediction will be human-interpretable in that users can identify features that determine the prediction of a store location's category.

2. An unsupervised machine learning pipeline based on data of US Subway locations that clusters locations by similar features.

3. An unsupervised machine learning pipeline based on data of Canadian Subway locations that clusters locations by similar features.

The two unsupervised machine learning pipelines will also have human-interpretable results, including ways to identify similar features that caused different locations to be clustered together.

The final product presented to Sitewise Analytics will be a GitHub repository containing the code for the machine learning data pipelines. Moreover, documentation of the scripts that can be run to reproduce the results will be included.

Data Summary

We received five datasets for each of the three popular chain restaurants: Smoothie King, Subway Canada, and Subway US. The datasets consist of CSV files for demographic, point of interest, store-specific data, competition sister store data, and trade area, where each row represents a single store location and the columns represent the variables/features of that store. All features in the demographic, point of interest, competition

sister store, and trade area files are numeric, whereas the store-specific data files contain categorical features such as state and market size.

For Smoothie King, there are over 1000 features combined for 796 stores.

For Subway US, there are over 1000 features combined for approximately 14,000 stores.

For Subway Canada, there are around 100 features combined for around 1,800 stores.

To enhance the understanding of the behavior of the data, we select some significant features and create visualizations to depict their distribution across categories. More information can be found in the `data` directory of the repository, given that it is infeasible to describe more features here.
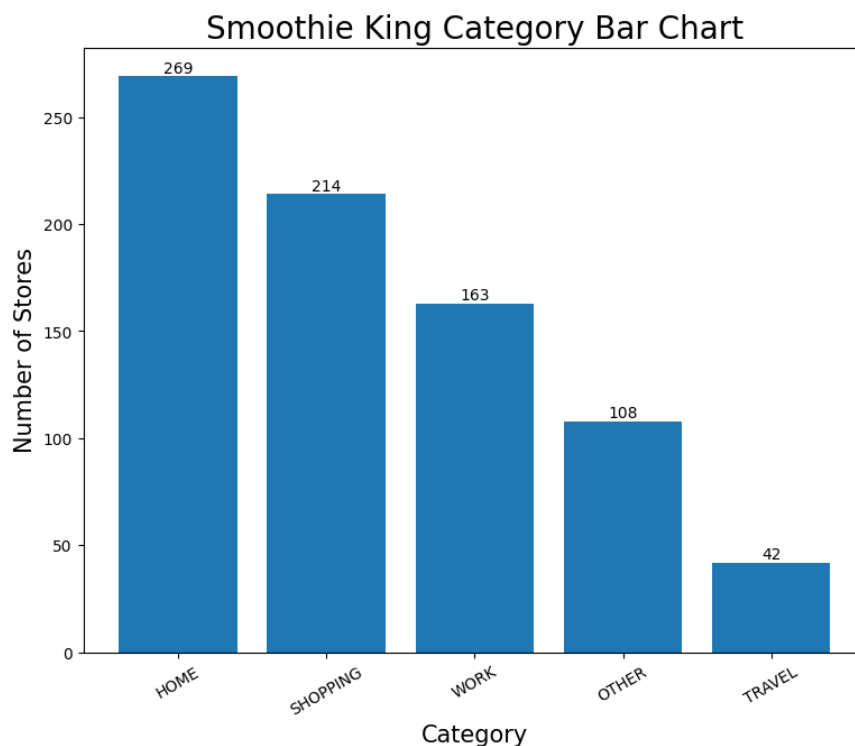
Understanding the Data

Figure 1: Number of stores in each category for Smoothie King US locations. The distribution is unbalanced.

Figure 2: The highest portion among categories are "Very Large Metro". "Home" has a high percentage of "Very Large Metro". "Travel" has a high percentage of "Small City".
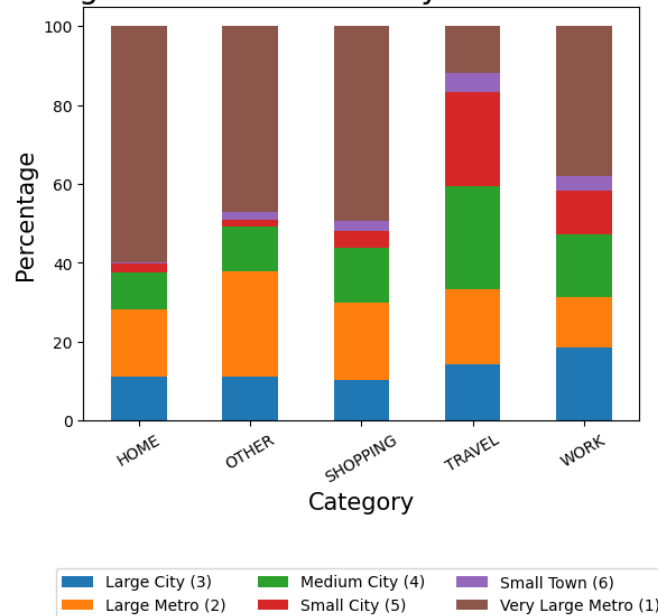

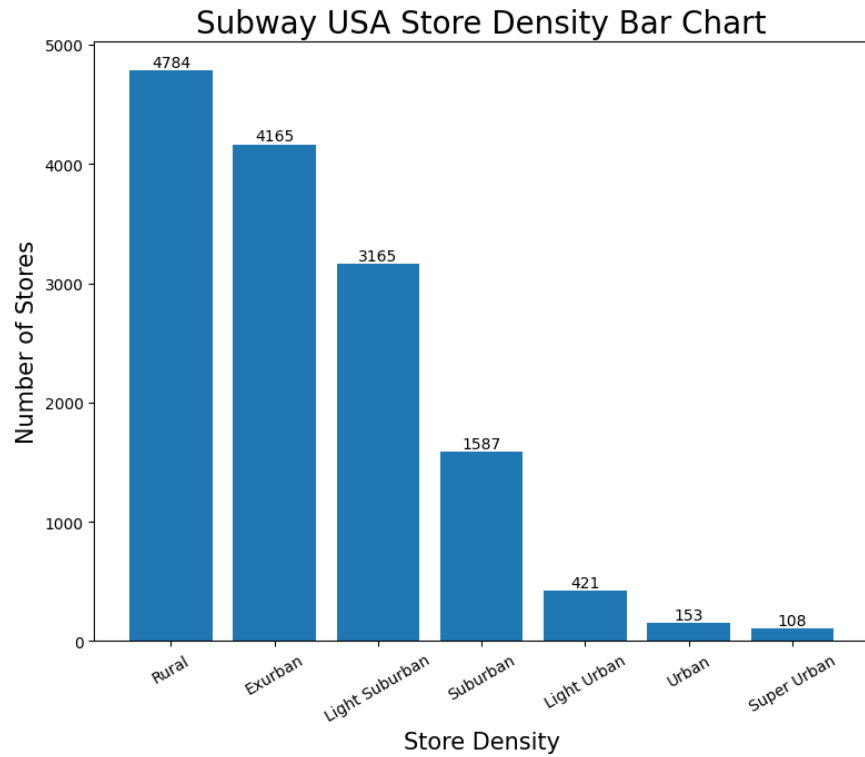Figure 3: Most of the stores are located in "Exurban" and "Light Suburban".

Figure 4: Rural has the highest count.

Figure 5: Very Large Metro has the highest count with 5190 stores, and other categories have similar counts.

Figure 6: The most common market size is Very Large Metro with 577 stores.

Figure 7: Very Large Metro has the highest count with 577 stores, and the other categories have similar counts.

Data Science Techniques

Based on the data described above with the important insights highlighted from some important features, along with the problems to be resolved for each restaurant chain which were mentioned in the *Introduction* section, the initial approaches and the corresponding evaluation metrics to be used are as followings:

1. For Smoothie King, many features are highly correlated to one another. Therefore, it is important to start with the feature selection and dimensionality reduction process. Other than manually analyzing the data columns with correlation scores, we are planning to use Principal Component Analysis (PCA) to reduce the dimension and Recursive Feature Elimination method to select the most important features.

   With the data and the supervised classification objective, we propose to start with a Logistic Regression model as our baseline. We will also explore other models such as a simple Random Forest Tree regression model, as well as an LGBMClassifier as it effectively trains high-dimensional datasets and usually leads to relatively high accuracy for the multi-class problems.

   With the labeled data, the models will be evaluated by its accuracy score, with the target being 80% accuracy as outlined by Sitewise. The score would indicate if the model is capable of detecting the most important traffic driving factor to a store. The result can be interpreted with SHAP (SHapley Additive exPlanations) plots where we can visualize which features drive the decision of a particular category assignment for a certain store.


Figure 8: Flow chart of Smoothie King

2. For Subway US, the main goal is to cluster the stores into different clusters where stores share similar features within the same cluster. We propose starting with a DBSCAN model as a baseline and evaluate the clusters empirically. Afterwards, we will perform Hierarchical clustering and apply PCA for dimensionality reduction and compare the results. The rationale for using Hierarchical clustering is that we do not have to pre-define the number of clusters and it also allows unbalanced cluster sizes.

   Since both Smoothie King and Subway US are both in the US and share similar features in the dataset, they could share similar segmentations to a certain extent. To perform PCA, the most driven features from the supervised model for Smoothie King can be used as a reference to evaluate this PCA step where more matched important features could indicate a more proper PCA.

   During the clustering, all possible linkage criteria would be tested to evaluate how to find similarities between clusters. Finally, since ground truth labels are not known, evaluation can be done through the Silhouette Coefficient where a higher Silhouette Coefficient score suggests that a model defines well separated clusters.

   To interpret the results and empirically validate the clusters, a potential approach is to randomly select a sample of stores in the same cluster and visualize them on Google Maps/Sitewise Internal application to check if they share similar geographic attributes (near highway exits, dense residential area, etc.).

3. For Subway Canada, we will take a similar approach to Subway US and evaluate the result.

   If the performance is not ideal, consider removing the PCA step and run clustering again to evaluate the result again with the same procedure performed on the Subway US model.


Figure 9: Flow chart of Subway US and Canada

Timeline

Schedule		Objective
Week 1 (5/1)	EDA & Proposal	Understand the problem, perform initial EDA on the dataset, and propose potential models and approaches to each objective.
Week 2 (5/8)	Feature Selection	Explore a variety of methods to determine the most important features for each of the three datasets.
Week 3 (5/15)	Supervised Model	Train and test a supervised classification model on the labeled dataset, as well as a list of the most important features as major indicators.
Week 4 (5/22)	Unsupervised Model (part I)	Train an unsupervised clustering model on one of the unlabeled datasets and apply several potential evaluation metrics.
Week 5 (5/29)	Unsupervised Model (part II)	Train an unsupervised clustering model on one of the unlabeled datasets and apply several potential evaluation metrics.
Week 6 (6/5)	Models Tuning & Parameter Optimization	Perform parameter tuning and optimization on the three models.
Week 7 (6/12)	Final Presentation	Present the final models as well as the list of important indicators from the list of features. Address potential directions and approaches for further study.
Week 8 (6/19)	Final Product & Reflection	Draft submission for the final product and report and iterate on feedback before final submission.